


# Deadlocks

## Deadlock illustration




- Bridge is one lane and viewed as resource
- Deadlocks can be resolved if one car backs up
- Several cars may need to back up

## Deadlock characterization: the 4 conditions

- Mutual exclusion
  - At least one non-sharable resource
  - Additional requests for resource delayed
- Hold and wait
  - Process holds at least one resource
  - Waiting to acquire resources held by other processes
- No preemption
  - Resources only released by a process after completing its task
- Circular wait
  - Set of processes  $\{P_0, P_1, \dots, P_n\}$
  - $P_0$  waiting for  $P_1$ ,  $P_1$  waiting for  $P_2$ , ...,  $P_{n-1}$  waiting for  $P_n$
  - $P_n$  waiting for  $P_0$

## Deadlock detection through a visual diagram




- System Resource Allocation Graphs
  - Processes (circle)
  - Resource (square, dots)
  - Directed Edges
    - Request
      - $P_i \rightarrow R_j$
    - Assignment
      - $R_j \rightarrow P_i$

## Methods for handling deadlocks

- Ensure system cannot have deadlocks (prevention or avoidance)
  - Deadlock Prevention
    - Eliminate at least one condition
      - Disadvantages
        - Low resource utilization
        - Reduced throughput
    - Prevent Mutual Exclusion
      - Needed for non-shareable resources but not for shareable ones
      - Cannot prevent mutex for all resources: some intrinsically non-shareable
    - Prevent Hold and Wait
      - Guarantee that process does not hold a resource when requesting another
      - Example: allocate all resource before execution
    - Prevent no Preemption

- One possibility: when requested not granted, processes must release all resources it is holding
  - Restarted only if it regains old and new resources
- Prevent Circular Wait
  - Impose ordering of resources
  - Require process to request resource in increasing number
    - If process holds  $R_j$ , it can request another  $R_j > R_i$
    - If process holds  $R_j$  and requests  $R_i$ , it must release  $R_j$
- Deadlock Avoidance
  - Needs information on how processes will request resources
  - Safe state
    - May allocate resources (up to maximum) in some order with no deadlock
    - Over-allocation: system unsafe
  - Includes *claim edge* (future request) denoted by  $P \rightarrow R$
  - Request granted only if no cycle is formed
  - Single instance resources only



- Multiple resource instances: Banker's algorithm
  - Ensure bank does not over-allocate cash
  - Determines if allocating requested resources maintains safe state
  - Less efficient than system = resource allocation graph
- Allow deadlocks then recover
- Ignore deadlocks (ostrich algorithm)
- Deadlock Detection
  - Algorithm to check for deadlocks
  - Algorithm to recover from deadlock
  - Detection algorithm depends on two factors:
    - How often is a deadlock likely to occur?
    - How many processes are affected by the deadlock?
  - Single instance resource
    - Use Wait-For graph
  - Several instances
    - Uses time-varying data structures found in Banker's Algorithm